

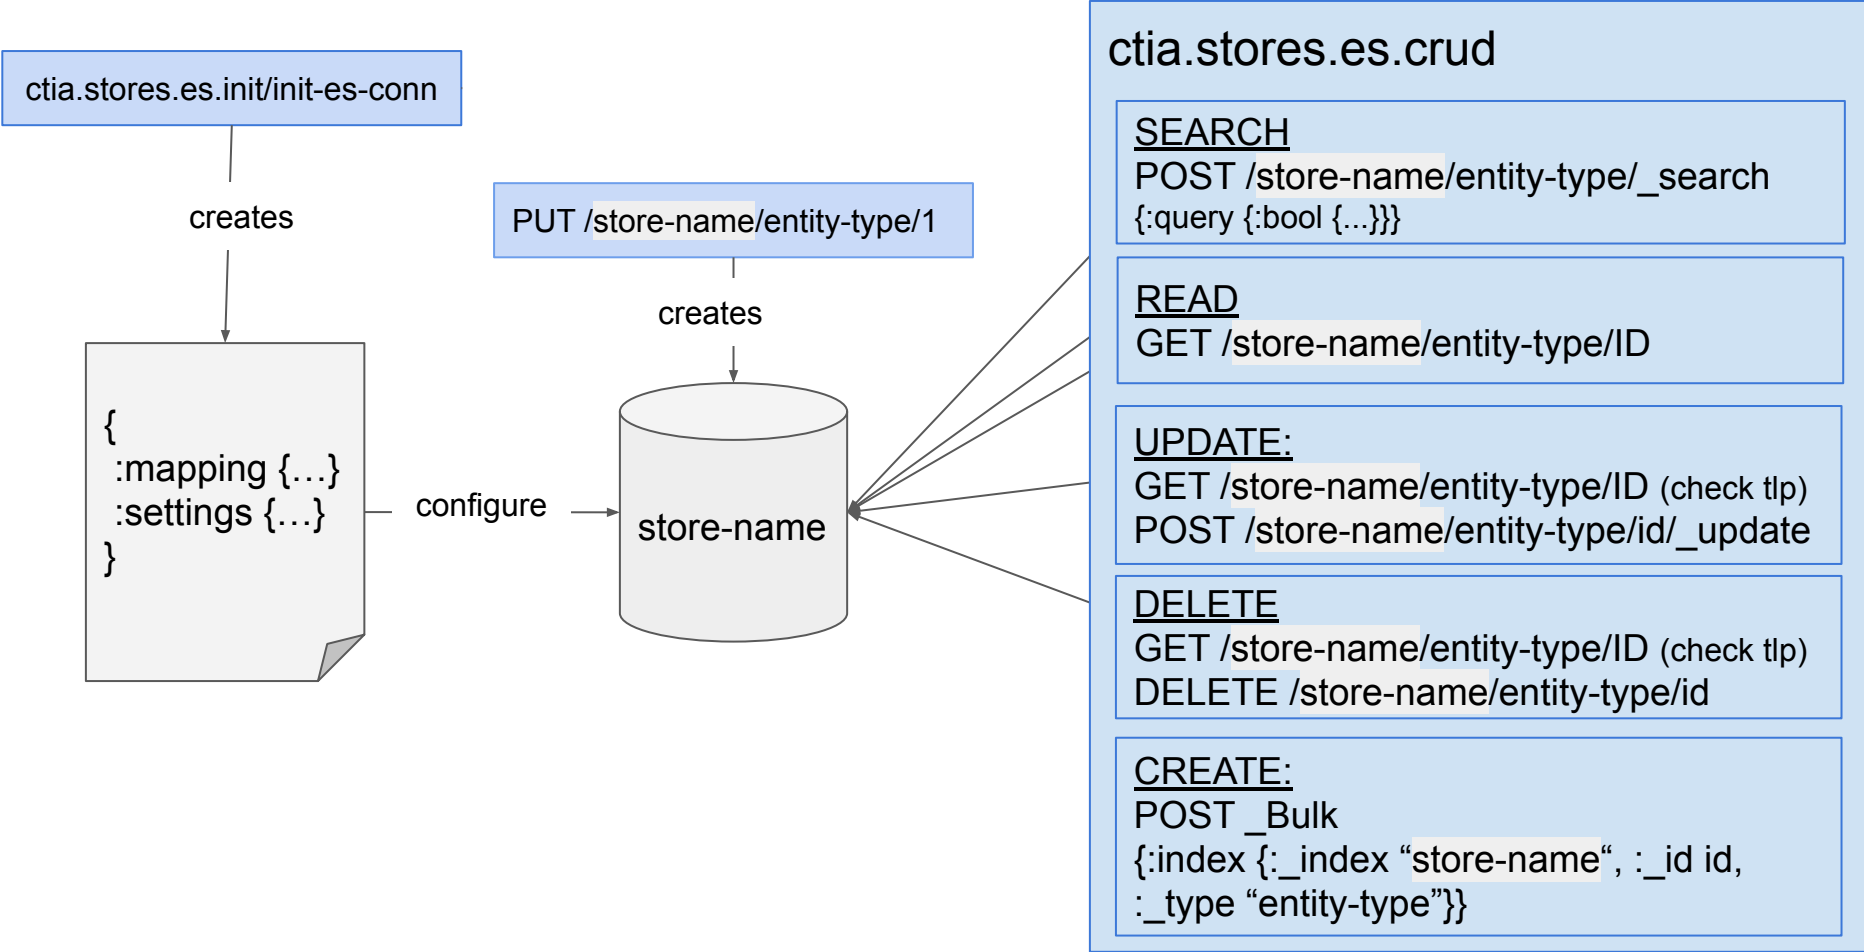
CTIA Elasticsearch stores

Managing big indices

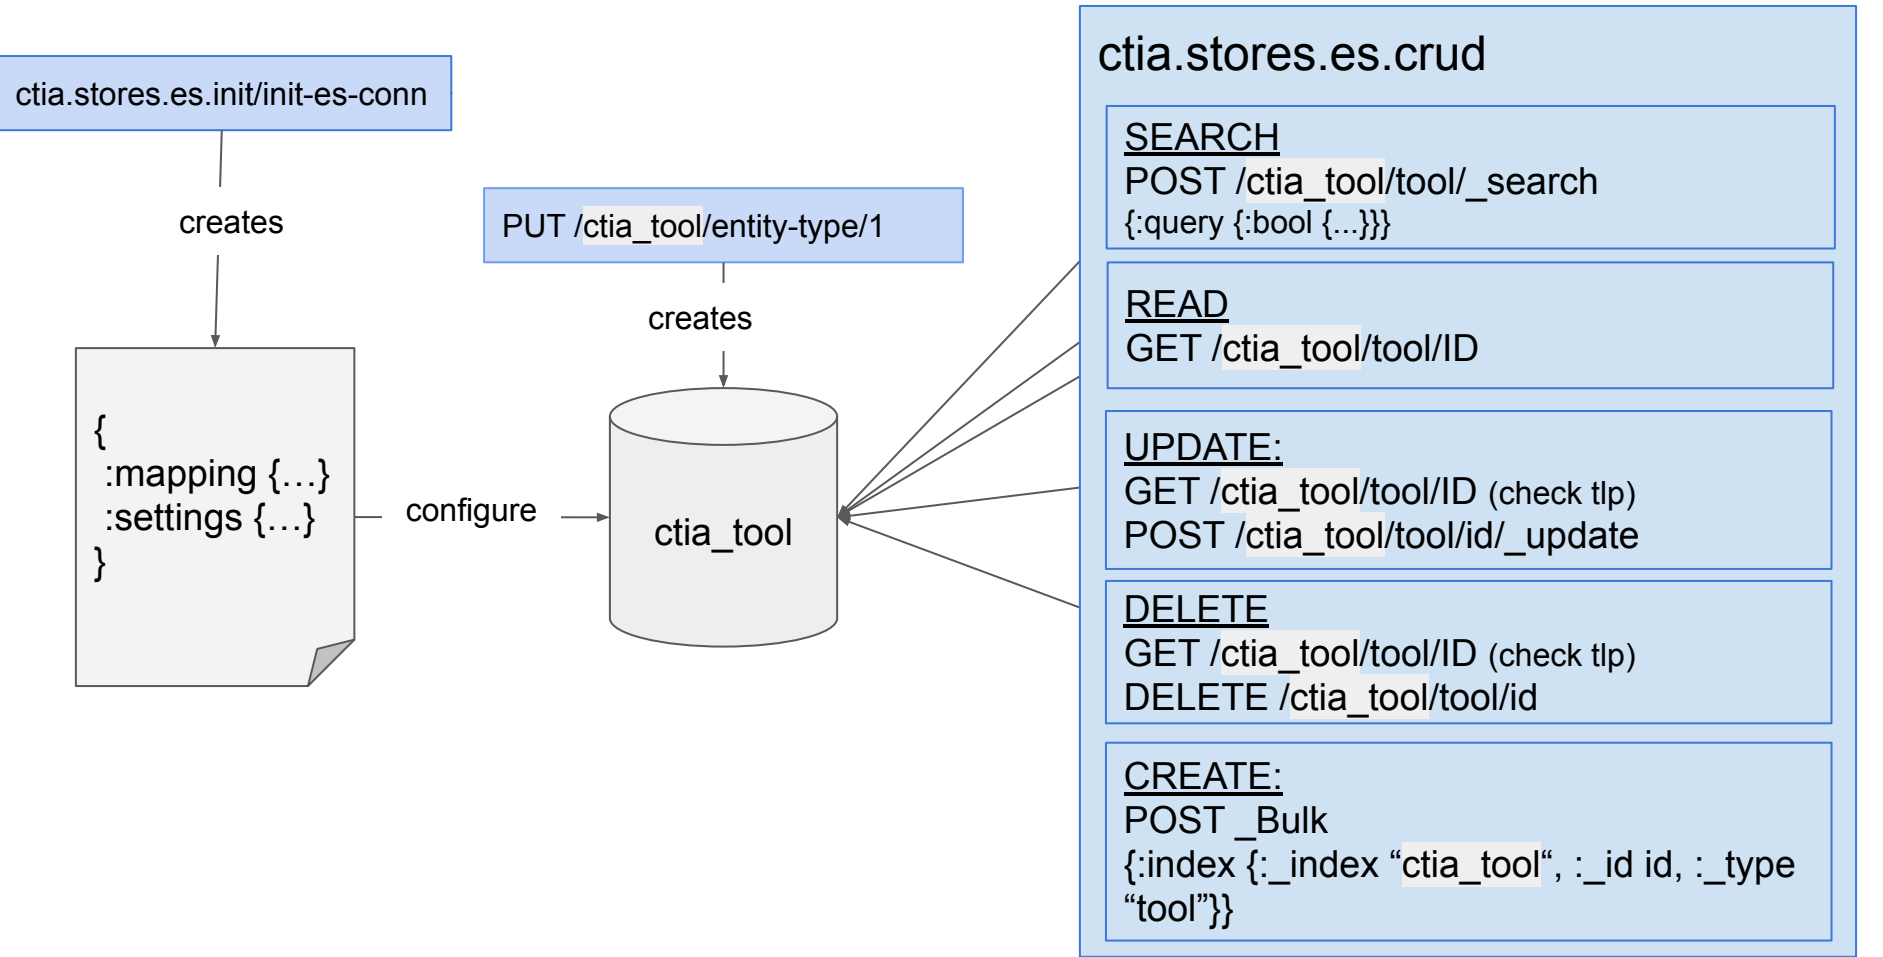
Unaliased stores: 1 store = 1 index

- Directly write on configured index name
- Initialization:
 - Create one template with proper index settings and type mapping.
 - Index is created with first document insertion, with proper configuration thanks to template
- CRUD operations and Search request directly target the index.
- + ⇒ simpler code
- - ⇒ limits horizontal scaling since number of shards cannot be increased.

unaliased stores: 1 store = 1 index



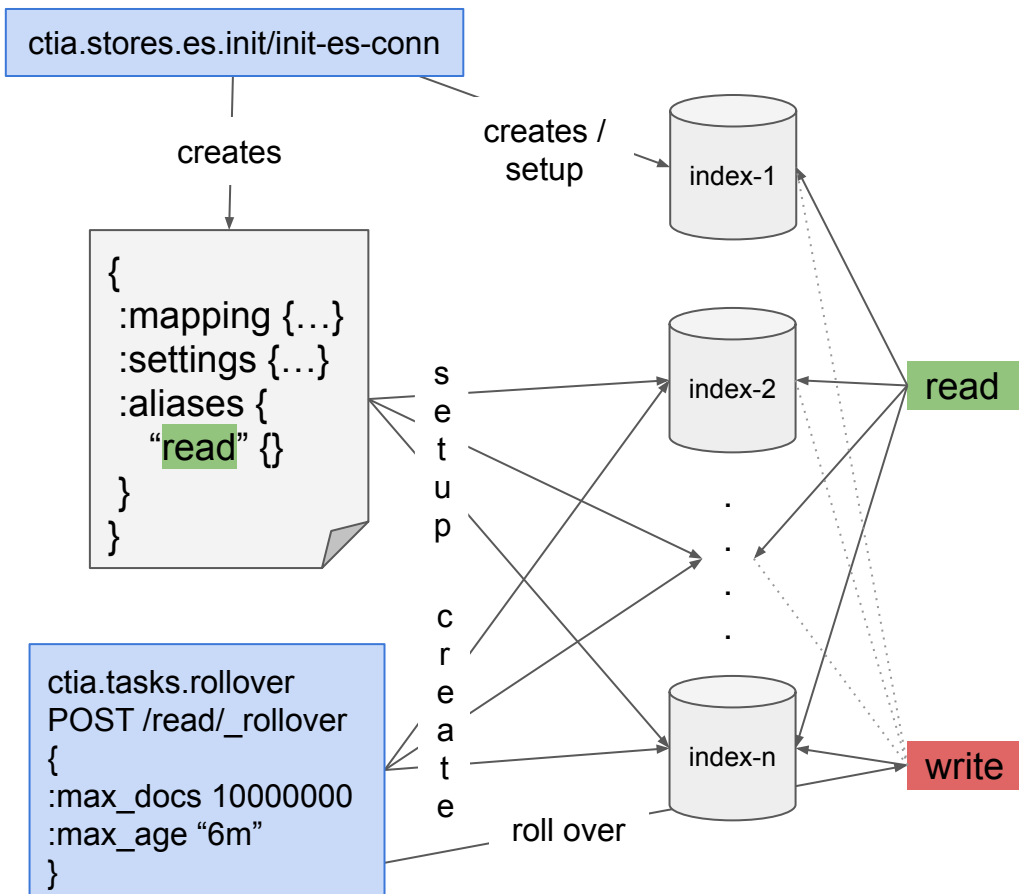
example: unaliased tool store



Aliased stores: 1 store / n indices, read/write aliases

- multiple indices, with *read* alias on all indices and *write* alias on most recent index. read alias is the store name, write alias is the store name suffixed by *-write*, index names are the store name suffixed by a number.
- Initialization: generates a template with a read alias and the first index with the write and read aliases.
- ES CRUD operations do not work on multiple indices:
 - GET is replaced by a `_search` request with an `ids` query on *read* alias
 - CREATE are done on the most recent index using *write* alias
 - UPDATE and DELETE require 2 steps:
 - 1) a `_search` with an `ids` query on *read* alias to retrieve the real index of the targeted document.
 - 2) update / delete request on that retrieved index.
- [_rollover](#) api is used to create a new index and to “roll over” *write* alias on that new index, when some rules (`max_docs` and / or `max_age`) are matched.
- + \Rightarrow unlimited horizontal scaling
- - \Rightarrow more complex code

Aliased stores: 1 store / n indices, read/write aliases



ctia.stores.es.crud

SEARCH

POST `/read/entity-type/_search`
{:query {:bool {...}}}

READ

POST `/read/entity-type/_search`
{"query": {"ids": {"values": [id]}}}

UPDATE:

POST `/read/entity-type/_search`
{"query": {"ids": {"values": [id]}}}
POST `/index-%d/entity-type/id/_update`

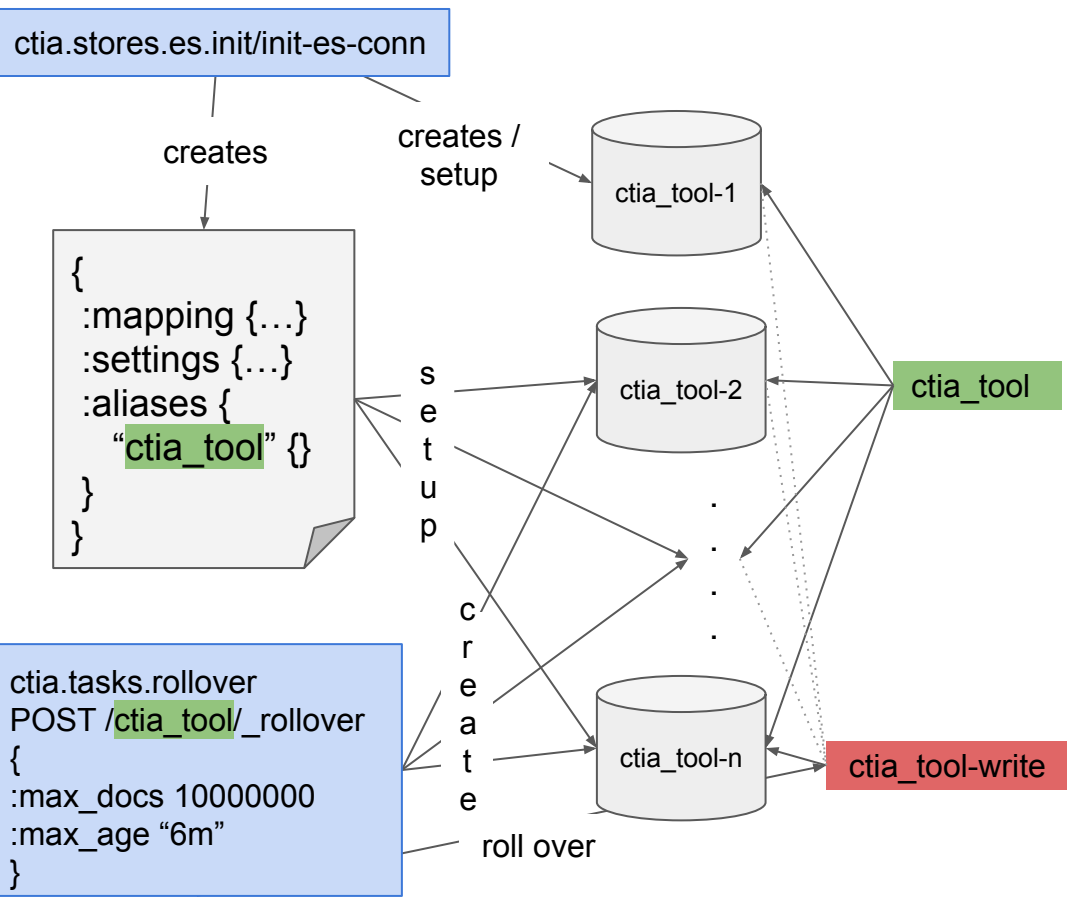
DELETE

POST `/read/entity-type/_search`
{"query": {"ids": {"values": [id]}}}
DELETE `/index-%d/entity-type/id`

CREATE

POST `_Bulk`
{:index {:_index "write", :_id id, :_type
"entity-type"}}

example: aliased tool store



ctia.stores.es.crud

SEARCH
POST /ctia-tool/entity-type/_search
{:query {:bool {...}}}

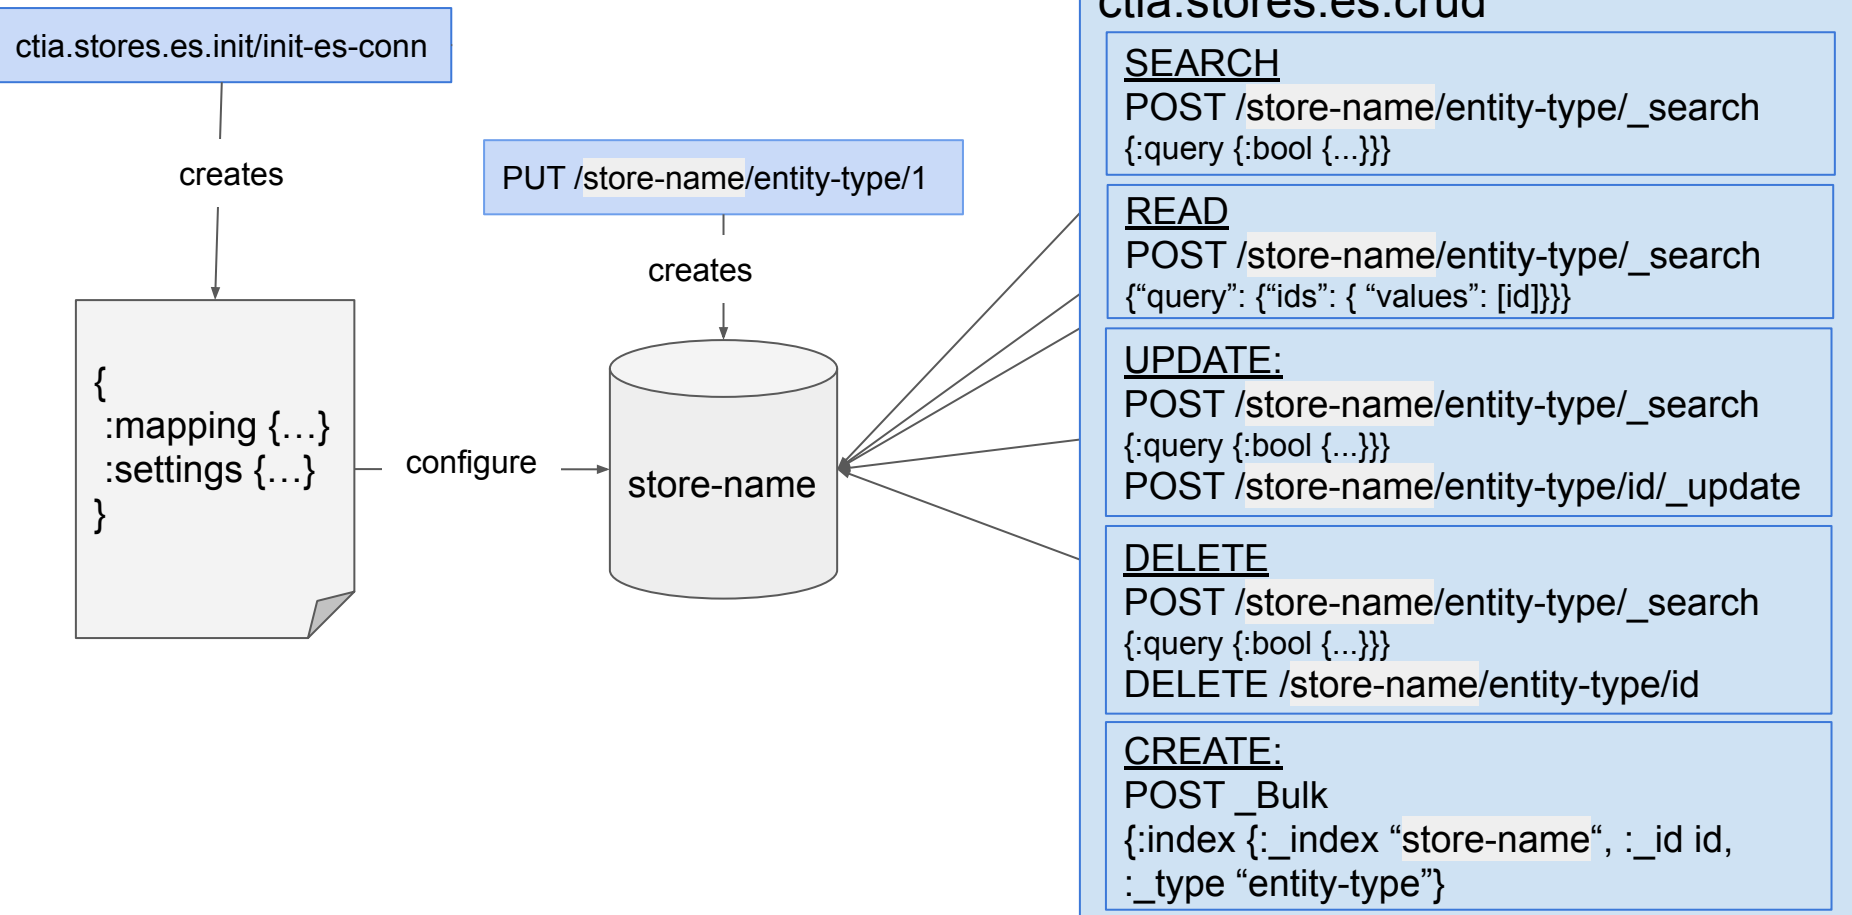
READ
POST /ctia-tool/entity-type/_search
{“query”: {“ids”: { “values”: [id]}}}

UPDATE:
POST /ctia-tool/entity-type/_search
{“query”: {“ids”: { “values”: [id]}}}
POST /ctia-tool-%d/entity-type/id/_update

DELETE
POST /ctia-tool/entity-type/_search
{“query”: {“ids”: { “values”: [id]}}}
DELETE /ctia-tool-%d/entity-type/id

CREATE
POST _Bulk
{:index {:_index “ctia_tool-write“, :_id id,
:_type “tool”}}

New unaliased stores: 1 store = 1 index + shared code



Migrating stores

- We have 4 main cases when migrating a store:
 - unaliased \Rightarrow unaliased
 - unaliased \Rightarrow aliased
 - aliased \Rightarrow unaliased
 - aliased \Rightarrow aliased
- In every case the problem is twofold:
 - Reading from source aliased or not.
 - Writing into target aliased or not.
- The migration can be started while CTIA still running, stopped and restarted if necessary. Thus it must detect updates and deletes in source store and apply it to target store:
 - We detect updates by sorting documents by modification date.
 - We detect deletes from event store.
- Moreover it must roll over target when an aliased index is wanted.

Migration: reading source.



Unaliased


Retrieve created and modified docs:

```
POST /store-name/entity-type/_search
{
  :sort [{:modified :desc},
        {:created :desc},
        {:uuid :desc}]
  :search_after: [..., ..., ...]
}
```

retrieve deleted docs from events

```
POST /event/event/_search
{
  :sort [{:modified :asc},
        {:created :asc},
        {:uuid :asc}]
  :search_after: [..., ..., ...]
}
```

Aliased



Retrieve created and modified docs:

```
POST /store-name/entity-type/_search
{
  :sort [{:modified :desc},
        {:created :desc},
        {:uuid :desc}]
  :search_after: [..., ..., ...]
}
```

retrieve deleted docs from events

```
POST /event/event/_search
{
  :sort [{:modified :asc},
        {:created :asc},
        {:uuid :asc}]
  :search_after: [..., ..., ...]
}
```

Migration: writing target



Unaliased

Bulk upsert for created and modified

POST /_bulk


```
{
  {:_index "v1.0.0-store-name", :mapping "entity-type", ...}
  {:_index "v1.0.0-store-name", :mapping "entity-type", ...}
  {:_index "v1.0.0-store-name", :mapping "entity-type", ...}
}
```

Delete_by_query deleted docs in source:

POST /v1.0.0-store-name/entity-type/_delete_by_query

```
{
  :query {
    :ids {:values [...]}
  }
}
```

Aliased



Bulk upsert for created and modified

POST /_bulk

```
{
  {:_index "store-name-write", :mapping "entity-type", ...}
  {:_index "index-1", :mapping "entity-type", ...}
  {:_index "index-3", :mapping "entity-type", ...}
}
```

Delete_by_query deleted docs in source
(works on multiple indices)

POST /store-name/entity-type/_delete_by_query

```
{
  :query {
    :ids {:values [...]}
  }
}
```

Rollover Target Store

When migrating, a *_rollover* request is sent after an insert of a batch when:

- the store is configured as *aliased*
- the size of current index exceed a multiple of *max_docs* by less than *batch-size*.
(*max_age* is not applicable).
- Code:

```
(s/defn rollover? [aliased? max_docs batch-size migrated-count]
  (and aliased?
        max_docs
        (> migrated-count max_docs)
        (<= (mod migrated-count max_docs)
             batch-size)))
```

Migration state to restart and trigger roll over

- The prefix used for target store, used to build the target index name from source name, *v{prefix}_source-name*.
- for each store it records the `started` and `completed` date. The start date is used to search deletes that occurred during migration in event store.
- State of reads in source and writes in target.
 - the `total` number of red documents in source and actually `migrated` in target (some document might be corrupted and thus not migrated, some documents might be red and migrated twice when updated during migration)
 - the `index` name of the source (the target index name is built using prefix).
 - the `search_after` value to use for next search in source

Migrating with limited downtime

<https://github.com/threatgrid/ctia/blob/master/migration.md>



lein run -m
ctia.task.migration.migrate-es-stores
--migrations 1.0.0
--prefix 1.1.0
--stores tool,malware
--batch-size 3000
--id migration-1
--confirm

- STOP CTIA
- lein run -m
ctia.task.migration.migrate-es-stores
--batch-size 3000
--id migration-1
--confirm
--restart

- Modify CTIA configuration with new store names
- Restart CTIA

Migrating with limited downtime but temporary inconsistency

<https://github.com/threatgrid/ctia/blob/master/migration.md>



```
lein run -m
ctia.task.migration.migrate-es-stores
--migrations 1.0.0
--prefix 1.1.0
--stores tool,malware
--batch-size 3000
--id migration-1
--confirm
```

- STOP CTIA
- lein run -m
ctia.task.migration.migrate-es-stores
--batch-size 3000
--id migration-1
--confirm
--restart

- Modify CTIA configuration with new store names
- Restart CTIA